

## Swift Local Refactoring

Xcode 9 includes a brand new refactoring engine. It can transform code locally within a single Swift source file, or globally, such as renaming a method or property that occurs in multiple files and even different languages. The logic behind local refactorings is implemented entirely in the compiler and SourceKit, and is now open source in the swift repository. Therefore, any Swift enthusiast can contribute refactoring actions to the language. This post discusses how a simple refactoring can be implemented and surfaced in Xcode.

### Kinds of Refactorings

A **local refactoring** occurs within the confines of a single file. Examples of local refactoring include *Extract Method* and *Extract Repeated Expression*. **Global refactorings**, which change code cross multiple files (such as *Global Rename*), currently require special coordination by Xcode and currently cannot be implemented on their own within the Swift codebase. This post focuses on local refactorings, which can be quite powerful in their own right.

A refactoring action is initiated by a user's cursor selection in the editor. According to how they are initialized, we categorize refactoring actions as cursor-based or range-based. **Cursor-based refactoring** has a refactoring target sufficiently specified by a cursor position in a Swift source file, such as rename refactoring. In contrast, **range-based refactoring** needs a start and end position to specify its target, such as Extract Method refactoring. To facilitate the implementation of these two categories, the Swift repository provides pre-analyzed results called `ResolvedCursorInfo` and `RangeInfo` to answer several common questions about a cursor position or a range in a Swift source file.

For instance, `ResolvedCursorInfo` can tell us whether a location in the source file points to the start of an expression and, if so, provide the corresponding compiler object of that expression. Alternatively, if the cursor points to a name, `ResolvedCursorInfo` gives us the declaration corresponding to that name. Similarly, `RangeInfo` encapsulates information about a given source range, such as whether the range has multiple entry or exit points.

To implement a new refactoring for Swift, we don't need to start from the raw representation of a cursor or a range position; instead, we can start with `ResolvedCursorInfo` and `RangeInfo` upon which a refactoring-specific analysis can be derived.

### Cursor-based Refactoring

Cursor-based Refactoring

Cursor-based refactoring is initiated by a cursor location in a Swift source file. Refactoring actions implement methods that the refactoring engine uses to display the available actions on the IDE and to perform the transformations.

Specifically, for displaying the available actions:

1. The user selects a location from the Xcode editor.
2. Xcode makes a request to sourcekitd to see what available refactoring actions exist for that location.
3. Each implemented refactoring action is queried with a `ResolvedCursorInfo` object to see if the action is applicable for that location.
4. The list of applicable actions is returned as response from sourcekitd and displayed to the user by Xcode.

When the user selects one of the available actions:

1. Xcode makes a request to sourcekitd to perform the selected action on the source location.
2. The specific refactoring action is queried with a `ResolvedCursorInfo` object, derived from the same location, to verify that the action is applicable.
3. The refactoring action is asked to perform the transformation with textual source edits.
4. The source edits are returned as response from sourcekitd and are applied by the Xcode editor.

To implement *String Localization* refactoring, we need to first declare this refactoring in the `RefactoringKinds.def` file with an entry like:

```
CURSOR_REFACTORING(LocalizeString, "Localize String", localize.string)
```

`CURSOR_REFACTORING` specifies that this refactoring is initialized at a cursor location and thus will use `ResolvedCursorInfo` in the implementation. The first field, `LocalizeString`, specifies the internal name of this refactoring in the Swift codebase. In this example, the class corresponding to this refactoring is named `RefactoringActionLocalizeString`. The string literal `"Localize String"` is the display name for this refactoring to be presented to users in the UI. Finally, `"localize.string"` is a stable key that identifies the refactoring action, which the Swift toolchain uses in communication with the source editor. This entry also allows the C++ compiler to generate the class stub for the String Localization refactoring and its callers. Therefore, we can focus on the implementation of the required functions.

After specifying this entry, we need to implement two functions to teach Xcode:

1. When it is appropriate to show the refactoring action.
2. What code change should be applied when a user invokes this refactoring action.

Both declarations are automatically generated from the aforementioned entry. To fulfill (1), we need to implement the `isApplicable` function of `RefactoringActionLocalizeString` in `Refactoring.cpp`, as below:

```
1 bool RefactoringActionLocalizeString::
2 isApplicable(ResolvedCursorInfo CursorInfo) {
3     if (CursorInfo.Kind == CursorInfoKind::ExprStart) {
```

```

4     if (auto *Literal = dyn_cast<StringLiteralExpr>(CursorInfo.TrailingExpr) {
5         return !Literal->hasInterpolation(); // Not real API.
6     }
7 }
8 }

```

Taking a `ResolvedCursorInfo` object as input, it's almost trivial to check when to populate the available refactoring menu with “localize string”. In this case, checking that the cursor points to the start of an expression (Line 3), and the expression is a string literal (Line 4) without interpolation (Line 5) is sufficient.

Next, we need to implement how the code under the cursor should be changed if the refactoring action is applied. To do this, we have to implement the `performChange` method of `RefactoringActionLocalizeString`. In the implementation of `performChange`, we can access the same `ResolvedCursorInfo` object that `isApplicable` received.

```

1 bool RefactoringActionLocalizeString::
2 performChange() {
3     EditConsumer.insert(SM, Cursor.TrailingExpr->getStartLoc(), "NSLocalizedString(");
4     EditConsumer.insertAfter(SM, Cursor.TrailingExpr->getEndLoc(), ", comment: \"\")");
5     return false; // Return true if code change aborted.
6 }

```

Still using String Localization as an example, the `performChange` function is fairly straightforward to implement. In the function body, we can use `EditConsumer` to issue textual edits around the expression pointed by the cursor with the appropriate Foundation API calls, as Lines 3 and 4 illustrate.

## Range-based Refactoring

### Range-based Refactoring

As the above figure shows, range-based refactoring is initiated by selecting a continuous range of code in a Swift source file. Taking the implementation of the *Extract Expression* refactoring as an example, we first need to declare the following item in `RefactoringKinds.def`.

```
RANGE_REFACTORING(ExtractExpr, "Extract Expression", extract.expr)
```

This entry declares that the *Extract Expression* refactoring is initiated by a range selection, named internally as `ExtractExpr`, using “`Extract Expression`” as display name, and with a stable key of “`extract.expr`” for service communication purposes.

To teach Xcode when this refactoring should be available, we also need to implement `isApplicable` for this refactoring in `Refactoring.cpp`, with the slight difference that the input is a `RangeInfo` instead of a `ResolvedCursorInfo`.

```

1 bool RefactoringActionExtractExpr::
2 isApplicable(ResolvedRangeInfo Info) {
3     if (Info.Kind != RangeKind::SingleExpression)
4         return false;
5     auto Ty = Info.getType();
6     if (Ty.isNull() || Ty.hasError())
7         return false;
8     ...
9     return true;
10 }

```

Though a little more complex than its counterpart in the aforementioned String Localization refactoring, this implementation is self-explaining too. Lines 3 to 4 check the kind of the given range, which has to be a single expression to proceed with the extraction. Lines 5 to 7 ensure the extracted expression has a well-formed type. Further conditions that need to be checked are omitted in the example for now. Interested readers can refer to `Refactoring.cpp` for more details. For the code change part, we can use the same `RangeInfo` instance to emit textual edits:

```

1 bool RefactoringActionExtractExprBase::performChange() {
2     llvm::SmallString<64> DeclBuffer;
3     llvm::raw_svector_ostream OS(DeclBuffer);
4     OS << tok::kw_let << " ";
5     OS << PreferredName;
6     OS << TyBuffer.str() << " = " << RangeInfo.ContentRange.str() << "\n";
7     Expr *E = RangeInfo.ContainedNodes[0].get<Expr*>();
8     EditConsumer.insert(SM, InsertLoc, DeclBuffer.str());
9     EditConsumer.insert(SM,
10                         Lexer::getCharSourceRangeFromSourceRange(SM, E->getSourceRange()),
11                         PreferredName)
12     return false; // Return true if code change aborted.
13 }

```

Lines 2 to 6 construct the declaration of a local variable with the initialized value of the expression under extraction, e.g. `let extractedExpr = foo()`. Line 8 inserts the declaration at the proper source location in the local context, and Line 9 replaces the original occurrence of the expression with a reference to the newly declared variable. As demonstrated by the code example, within the function body of `performChange`, we can access not only the original `RangeInfo` for the user's selection, but also other important utilities such as the edit consumer and source manager, making the implementation more convenient.

## Diagnostics

A refactoring action may need to be aborted during automated code change for various reasons. When this happens, a refactoring implementation can

communicate via diagnostics the cause of such failures to the user. Refactoring diagnostics employ the same mechanism as the compiler itself. Taking rename refactoring as an example, we would like to issue an error message if the given new name is an invalid Swift identifier. To do so, we first need to declare the following entry for the diagnostics in `DiagnosticsRefactoring.def`.

```
ERROR(invalid_name, none, "'%0' is not a valid name", (StringRef))
```

After declaring it, we can use the diagnostic in either `isApplicable` or `performChange`. For *Local Rename* refactoring, emitting the diagnostic in `Refactoring.cpp` would look something like:

```
1 bool RefactoringActionLocalRename::performChange() {
2     ...
3     if (!DeclNameViewer(PreferredName).isValid()) {
4         DiagEngine.diagnose(SourceLoc(), diag::invalid_name, PreferredName);
5         return true; // Return true if code change aborted.
6     }
7     ...
8 }
```

## Testing

Corresponding to the two steps in implementing a new refactoring action, we need to test that:

1. The contextually available refactorings are populated properly.
2. The automated code change updates the user’s codebase correctly.

These two parts are both tested using the `swift-refactor` command line utility which is built alongside the compiler.

### Contextual Refactoring Test

```
1 func foo() {
2     print("Hello World!")
3 }
4 // RUN: %refactor -source-filename %s -pos=2:14 | %FileCheck %s -check-prefix=CHECK-LOCALIZE-STRING
5 // CHECK-LOCALIZE-STRING: Localize String
```

Let’s again take String Localization as an example. The above code snippet is a test for contextual refactoring actions. Similar tests can be found in `test/refactoring/RefactoringKind/`.

Let’s take a look at the `RUN` line in more detail, starting with the use of the `%refactor` utility:

```
%refactor -source-filename %s -pos=2:14 | %FileCheck %s -check-prefix=CHECK-LOCALIZE-STRING
```

This line will dump the display names for all applicable refactorings when a user points the cursor to the string literal “Hello World!”. `%refactor` is an alias that

gets substituted by the test runner to give the full path to `swift-refactor` when the tests get run. `-pos` gives the cursor position where contextual refactoring actions should be pulled from. Since `String Localization` refactoring is cursor-based, specifying `-pos` alone will be sufficient. To test range-based refactorings, we need to specify `-end-pos` to indicate the end location of the refactoring target as well. All positions are in the format of `line:column`.

To make sure the output of the tool is the expected one, we use the `%FileCheck` utility:

```
%FileCheck %s -check-prefix=CHECK-LOCALIZE-STRING
```

This will check the output text from `%refactor` against all following lines with prefix `CHECK-LOCALIZE-STRING`. In this case, it will check whether the available refactorings include `Localize String`. In addition to testing that we show the right actions at the right cursor positions, we also need to test available refactorings are not wrongly populated in situations like string literals with interpolation.

**Code Transformation Test** We should also test that when applying the refactoring, the automated code change matches our expectations. As a preparation, we need to teach `swift-refactor` a refactoring kind flag to specify the action we are testing with. To achieve this, the following entry is added in `swift-refactor.cpp`:

```
clEnumValN(RefactoringKind::LocalizeString, "localize-string", "Perform String Localizati
```

With such an entry, `swift-refactor` can test the code transformation part of `String Localization` specifically. A typical code transformation test consists of two parts:

1. The code snippet before refactoring.
2. The expected output after transformation.

The test performs the designated refactoring in (1) and compares the result with (2). It passes if the two are identical, otherwise the test fails.

```
1 func foo() {
2     print("Hello World!")
3 }
4 // RUN: rm -rf %t.result && mkdir -p %t.result
5 // RUN: %refactor -localize-string -source-filename %s -pos=2:14 > %t.result/localized.s
6 // RUN: diff -u %S/Inputs/localized.swift.expected %t.result/localized.swift

1 func foo() {
2     print(NSLocalizedString("Hello World!", comment: ""))
3 }
```

The above two code snippets comprise a meaningful code transformation test. Line 4 prepares a temporary source directory for the code resulting from the refactoring; using the newly added `-localize-string`, Line 5 performs the

refactoring code change at the start position of "Hello World!" and dumps the result to the temporary directory; finally, Line 6 compares the result with the expected output illustrated in the second code example.

## Integrating with Xcode

After implementing all of above pieces in the Swift codebase, we are ready to test/use the newly added refactoring in Xcode by integrating with a locally-built open source toolchain.

1. Run build-toolchain to build the open source toolchain locally.
2. Untar and copy the toolchain to `/Library/Developer/Toolchains/`.
3. Specify the local toolchain for Xcode's use via `Xcode->Toolchains`, like the following figure illustrates.

Specify Toolchain

## Potential Local Refactoring Ideas

This post just touches on some of the things that are now possible to implement in the new refactoring engine. If you are excited about extending the refactoring engine to implement additional transformations, Swift's issue database contains several ideas of refactoring transformations awaiting implementations.

For further help with implementing refactoring transformations, please feel free to ask questions on the Swift forums.